Practical 4: EDA

EDA

[14]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings as wr
wr.filterwarnings('ignore')
```

[2]:
```
df =pd.read_csv("winequality-red.csv")
```

[3]:
```
df
```

[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6

[4]:
```
df.shape
```

[4]:
```
(1599, 12)
```

[5]:
```
df.describe
```

[5]:
```
<bound method NDFrame.describe of       fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0               7.4             0.700         0.00             1.9      0.076
1               7.8             0.880         0.00             2.6      0.098
2               7.8             0.760         0.04             2.3      0.092
3              11.2             0.280         0.56             1.9      0.075
4               7.4             0.700         0.00             1.9      0.076
...             ...               ...          ...             ...        ...
1594            6.2             0.600         0.08             2.0      0.090
1595            5.9             0.550         0.10             2.2      0.062
1596            6.3             0.510         0.13             2.3      0.076
1597            5.9             0.645         0.12             2.0      0.075
1598            6.0             0.310         0.47             3.6      0.067
```

[6]:
```
df.isnull().sum()
```

[6]:
```
fixed acidity           0
volatile acidity        0
citric acid             0
residual sugar          0
chlorides               0
free sulfur dioxide     0
total sulfur dioxide    0
density                 0
pH                      0
sulphates               0
alcohol                 0
quality                 0
dtype: int64
```

[7]:
```
df.columns.tolist()
```

[7]:
```
['fixed acidity',
 'volatile acidity',
 'citric acid',
 'residual sugar',
 'chlorides',
 'free sulfur dioxide',
 'total sulfur dioxide',
 'density',
 'pH',
 'sulphates',
 'alcohol',
 'quality']
```

[8]:
```
df.nunique()
```

[8]:
```
fixed acidity           96
volatile acidity       143
citric acid             80
residual sugar          91
chlorides              153
free sulfur dioxide     60
total sulfur dioxide   144
density                436
pH                      89
sulphates               96
alcohol                 65
quality                  6
dtype: int64
```

[9]:
```
qc =df['quality'].value_counts()

plt.figure(figsize=(8,6))
plt.bar(qc.index,qc,color="red")
plt.title("Count plot of Quality")
plt.xlabel("Quality")
plt.ylabel("Count")
plt.show()
```

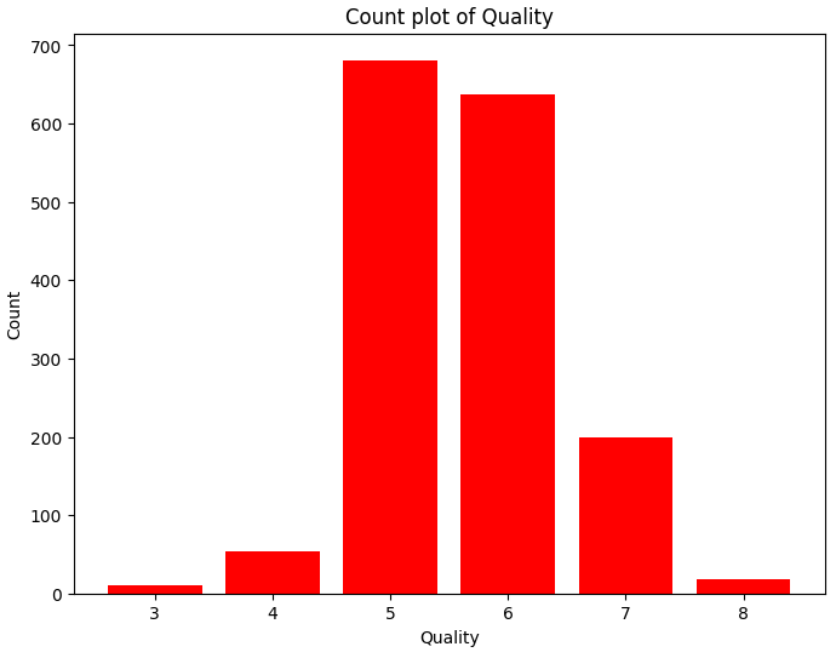

Kernel Density Plots

[11]:
```
sns.set_style("darkgrid")
numerical_columns = df.select_dtypes(include=["int64", "float64"]).columns

plt.figure(figsize=(14, len(numerical_columns) * 3))
for idx, feature in enumerate(numerical_columns, 1):
 plt.subplot(len(numerical_columns), 2, idx)
 sns.histplot(df[feature], kde=True)
 plt.title(f"{feature} | Skewness: {round(df[feature].skew(), 2)}")

plt.tight_layout()
plt.show()
```

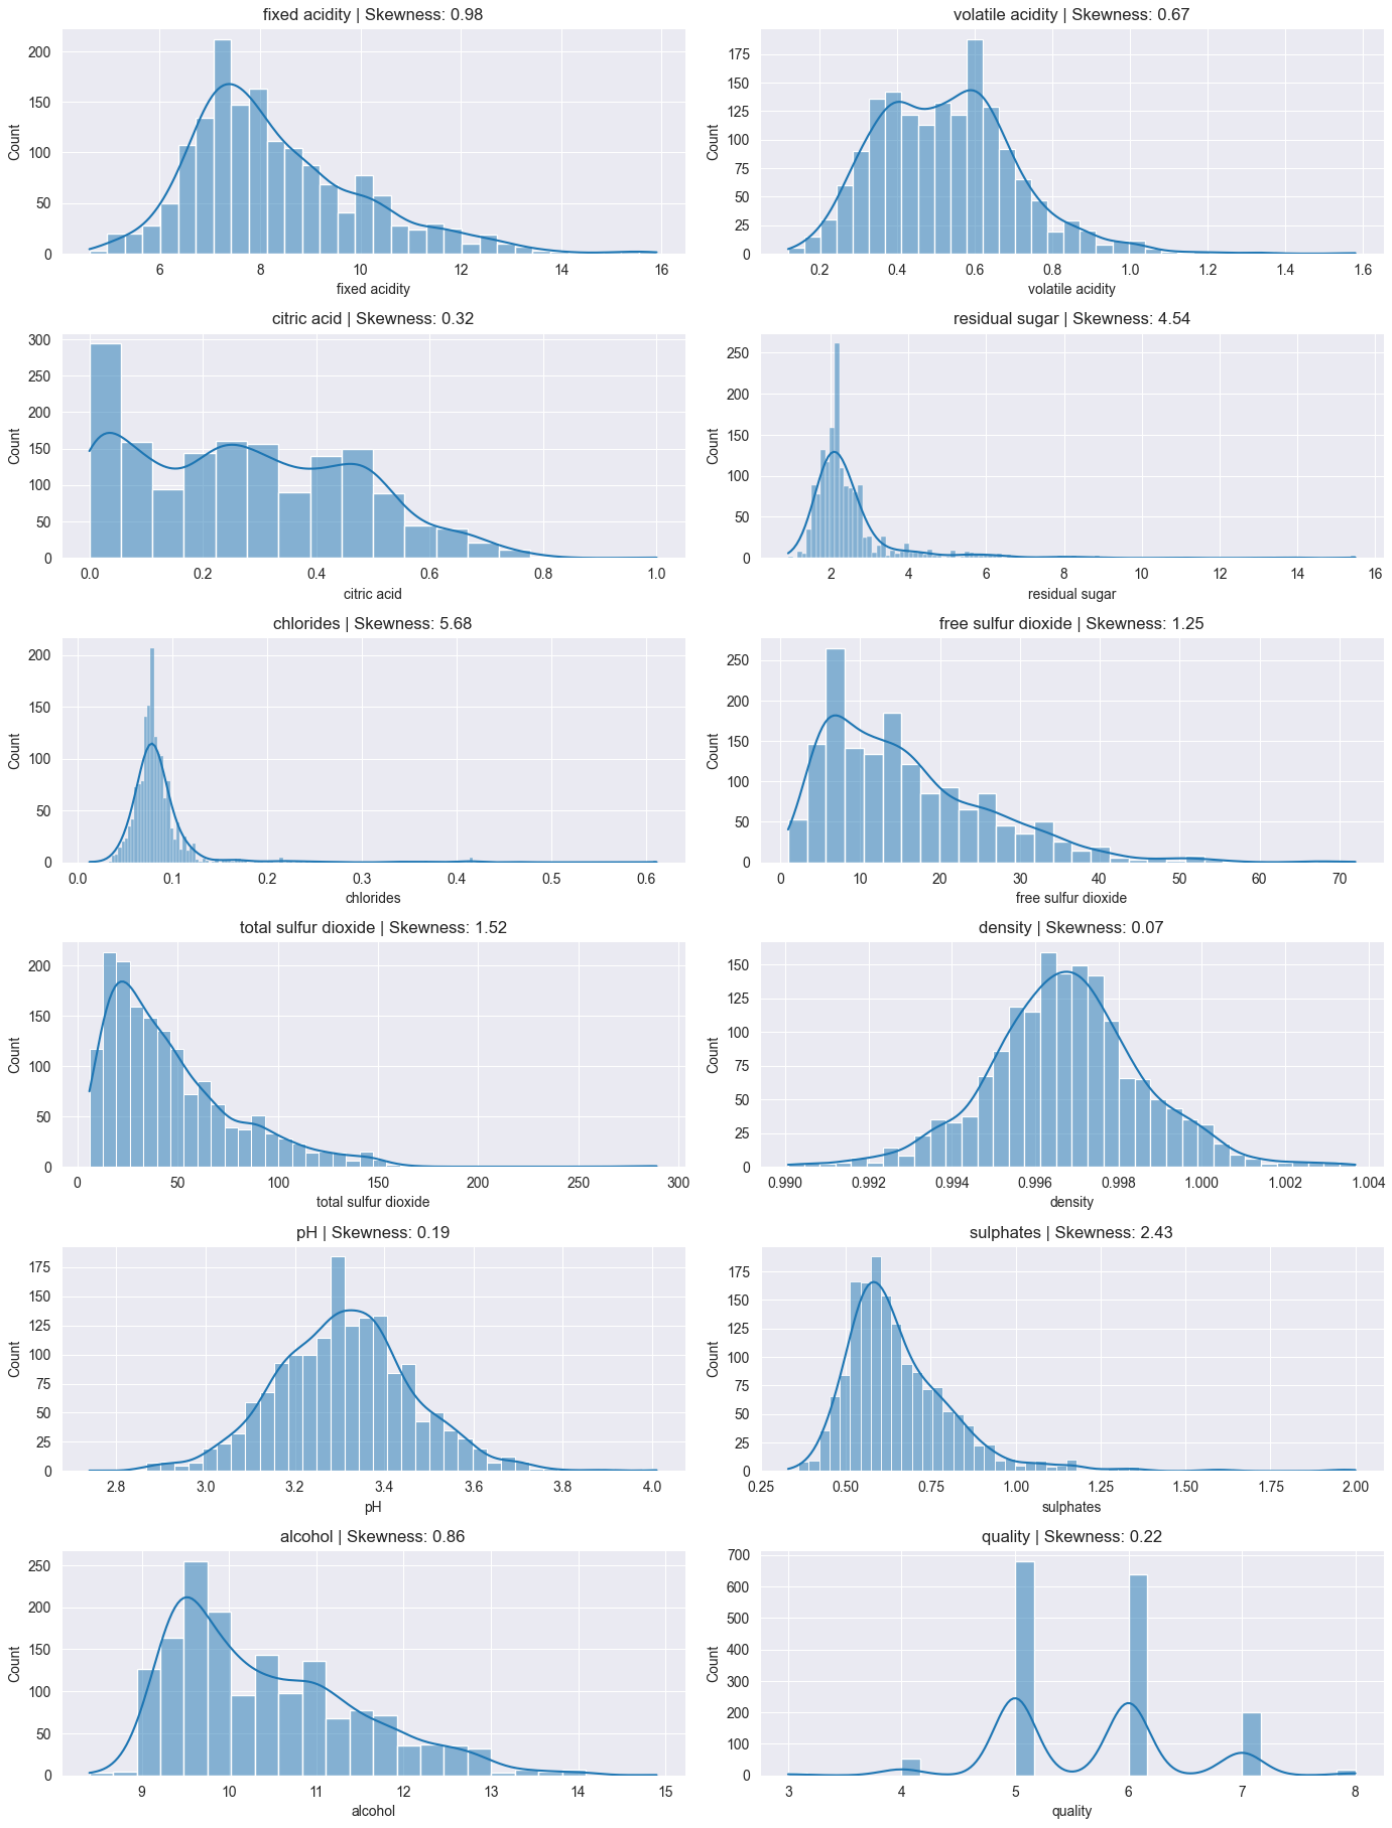
fixed acidity | Skewness: 0.98
Count
fixed acidity
volatile acidity | Skewness: 0.67
Count
volatile acidity
citric acid | Skewness: 0.32
Count
citric acid
residual sugar | Skewness: 4.54
Count
residual sugar
chlorides | Skewness: 5.68
Count
chlorides
free sulfur dioxide | Skewness: 1.25
Count
free sulfur dioxide
total sulfur dioxide | Skewness: 1.52
Count
total sulfur dioxide
density | Skewness: 0.07
Count
density
pH | Skewness: 0.19
Count
pH
sulphates | Skewness: 2.43
Count
sulphates
alcohol | Skewness: 0.86
Count
alcohol
quality | Skewness: 0.22
Count
quality

Swarm plot

[15]:
```
plt.figure(figsize=(10, 8))

sns.swarmplot(x="quality", y="alcohol", data=df, palette='viridis')
plt.title('Swarm Plot for Quality and Alcohol')
plt.xlabel('Quality')
plt.ylabel('Alcohol')
plt.show()
```

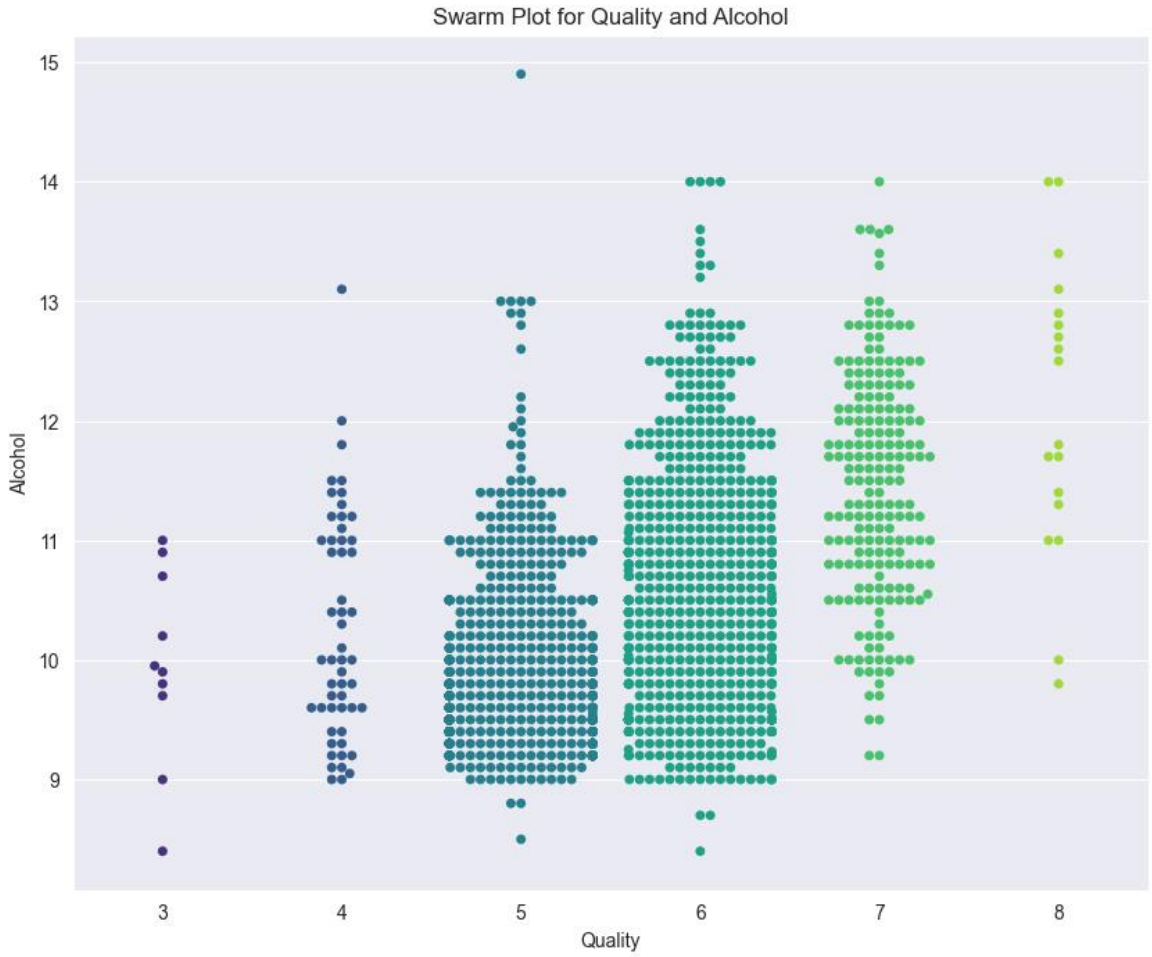

Pair Plots

[17]:
```
sns.set_palette("Pastel1")
plt.figure(figsize=(10, 6))
sns.pairplot(df)
plt.suptitle('Pair Plot for DataFrame')
plt.show()
```

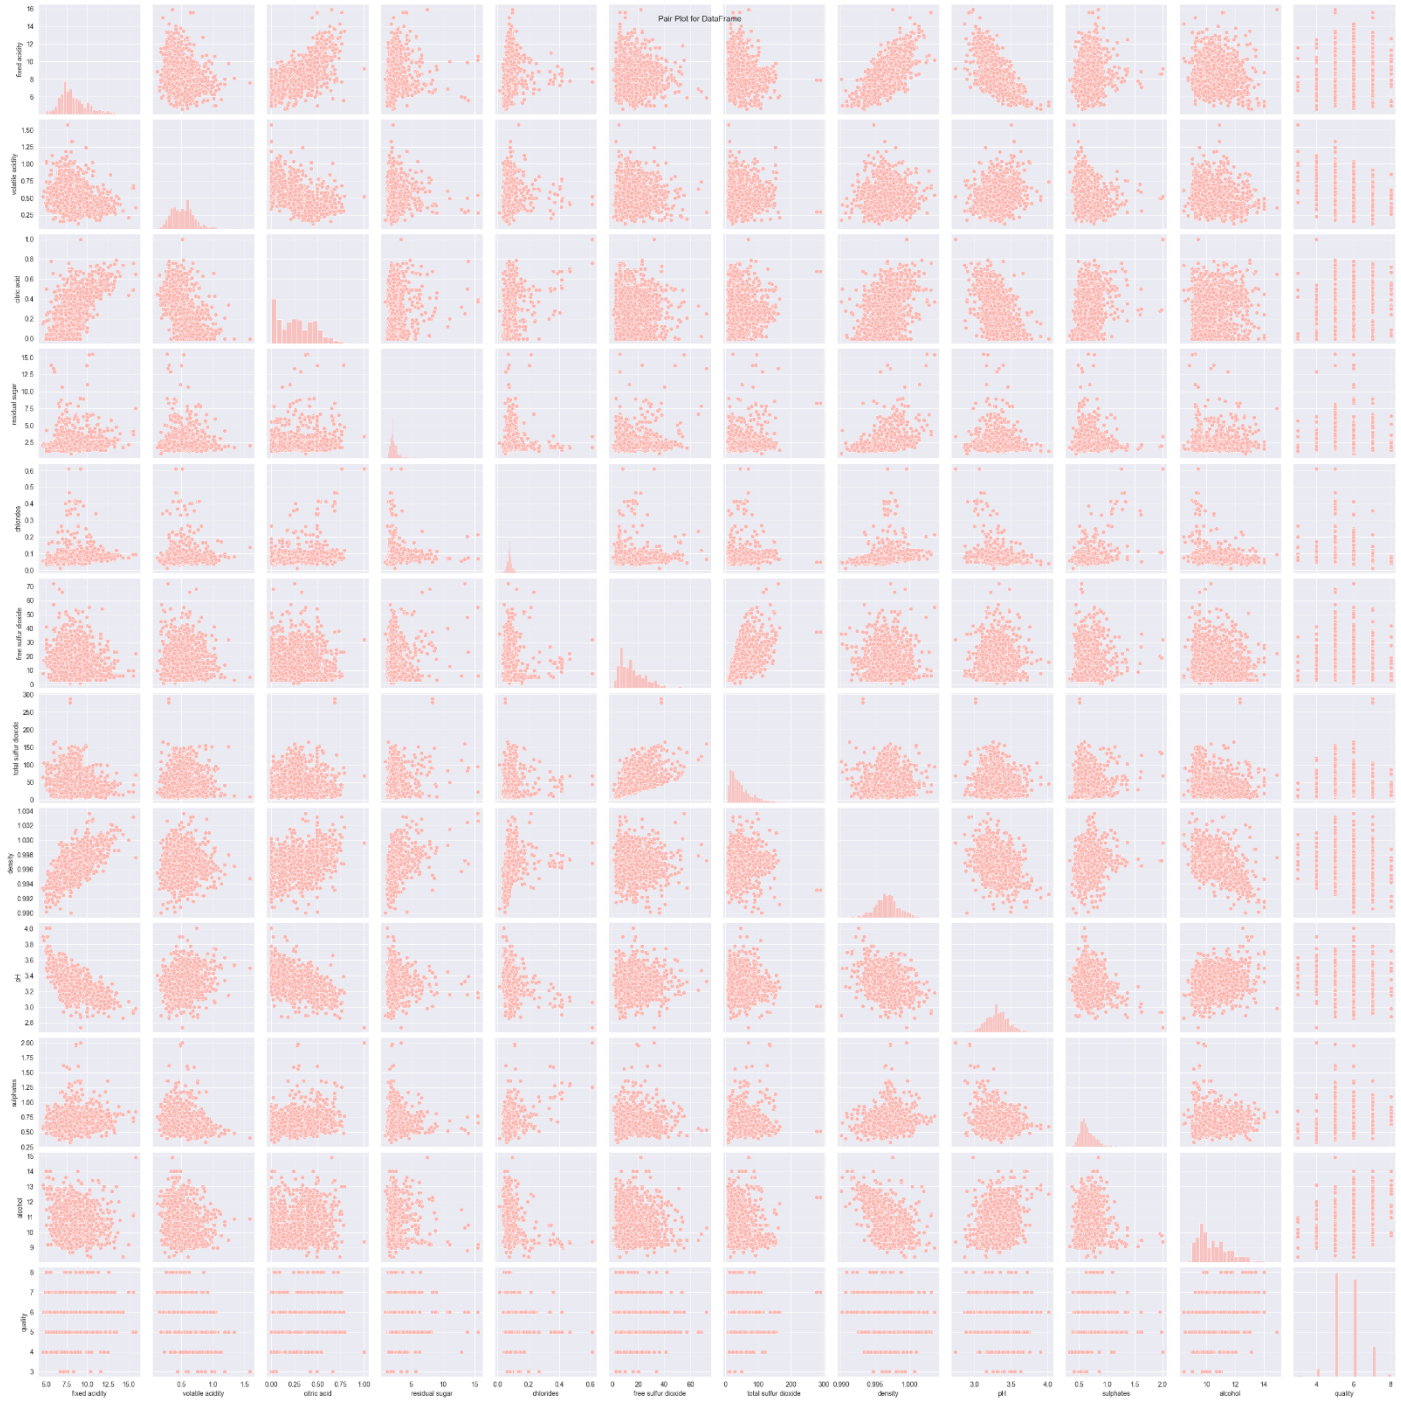

Violin plot

[18]:
```
df['quality'] = df['quality'].astype(str) # Convert 'quality' to categorical
plt.figure(figsize=(10, 8))
sns.violinplot(x="quality", y="alcohol", data=df, palette={
    '3': 'lightcoral', '4': 'lightblue', '5': 'lightgreen', '6': 'gold',
    '7': 'lightskyblue', '8': 'lightpink'},
    alpha=0.7)
plt.title('Violin Plot for Quality and Alcohol')
plt.xlabel('Quality')
plt.ylabel('Alcohol')
plt.show()
```

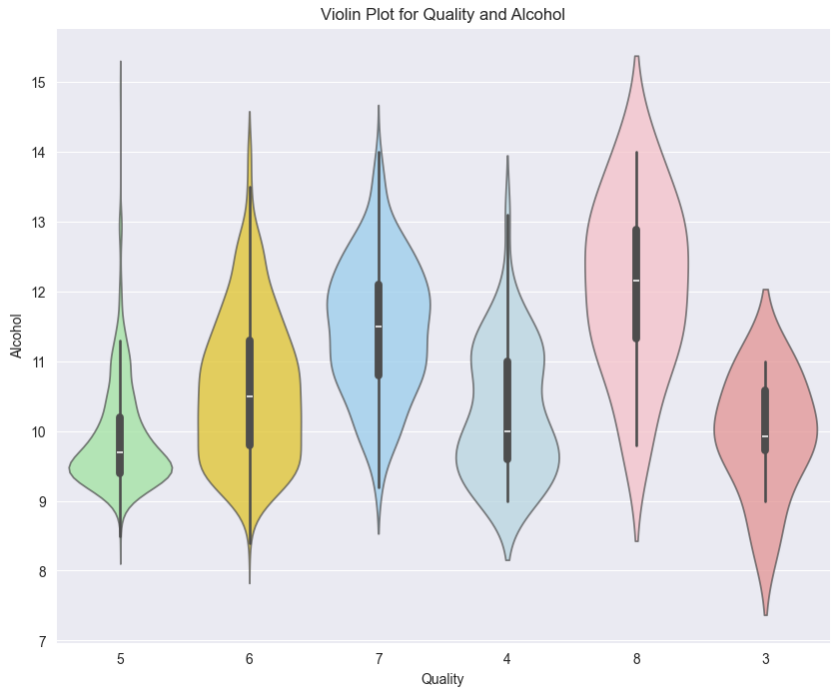

Box Plot

[20]:
```
sns.boxplot(x='quality', y='alcohol', data=df)
```

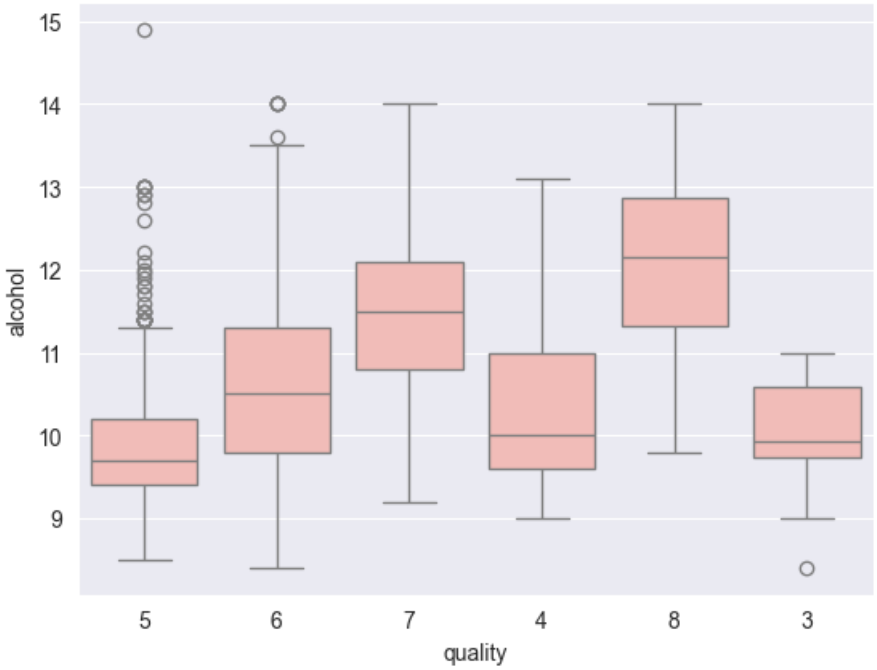

Correlation Matrix

[25]:
```
plt.figure(figsize=(15, 10))
sns.heatmap(df.corr(), annot=True, fmt='.2f', cmap='Blues', linewidths=2)
plt.title('Correlation Heatmap')
plt.show()
```